

Susie-relax with the “irrelevance of the null SNPs” property

1 SER-relax-prime

Now we consider another approximation to truly achieve the “irrelevance of the null SNPs” property. It leads to another model that I call ‘SER-relax-prime’.

Recall the SER-relax (conditional on $\gamma = 1$):

$$x \mid \beta, \gamma = 1 \sim \mathcal{N}\left(\beta \begin{pmatrix} 1 \\ \mu_1 \end{pmatrix}, \frac{1}{N} \begin{pmatrix} 1 & \mu_1^T \\ \mu_1 & \bar{R}_{-1,-1} + N\beta^2\tau^2 S_1 \end{pmatrix}\right) \quad (1)$$

Matthew + Gemini suggest omitting β^2 (replacing β^2 by 1) from the variance in the formula above. Similar to computation in PIP from the previous file, we have

$$\begin{aligned} p(x \mid \beta, \gamma = j) &= \frac{N^{J/2}}{(2\pi)^{J/2} |\bar{R}|^{1/2} (1 + N\tau^2)^{(J-1)/2}} \exp\left(-\frac{N}{2(1 + N\tau^2)} q\right) \\ &\times \exp\left(-\frac{N^2\tau^2}{2(1 + N\tau^2)} x_j^2\right) \exp\left(N\beta x_j - \frac{N}{2}\beta^2\right). \end{aligned} \quad (2)$$

Integrating out $\beta \sim \mathcal{N}(0, \sigma_0^2)$, we have

$$\begin{aligned} p(x \mid \gamma = j) &= \frac{N^{J/2}}{(2\pi)^{J/2} |\bar{R}|^{1/2} (1 + N\tau^2)^{(J-1)/2}} \exp\left(-\frac{N}{2(1 + N\tau^2)} q\right) \\ &\times \sqrt{\frac{1}{1 + N\sigma_0^2}} \times \exp\left(\left(\frac{N\sigma_0^2}{1 + N\sigma_0^2} - \frac{N\tau^2}{1 + N\tau^2}\right) \frac{Nx_j^2}{2}\right). \end{aligned} \quad (3)$$

Leading to

$$p(\gamma = j \mid x) \propto \exp\left(\left(\frac{N\sigma_0^2}{1 + N\sigma_0^2} - \frac{N\tau^2}{1 + N\tau^2}\right) \frac{Nx_j^2}{2}\right). \quad (4)$$

So we truly have the desired property. However, something is not right about the formula above: The scalar in the exponential can turn into negative if $\sigma_0^2 < \tau^2$ (so that a smaller Z -score leads to a bigger PIP?). It is because we are replacing β^2 by 1 in the variance of the likelihood of $p(x \mid \beta, \gamma = j)$. Maybe the approximation $\beta^2 \approx \sigma_0^2$ makes more sense.

SER-relax-prime. This model approximates β^2 in the variance of (1) by σ_0^2 . It leads to

$$x \mid \beta, \gamma = 1 \sim \mathcal{N}\left(\beta \begin{pmatrix} 1 \\ \mu_1 \end{pmatrix}, \frac{1}{N} \begin{pmatrix} 1 & \mu_1^T \\ \mu_1 & \bar{R}_{-1,-1} + N\sigma_0^2\tau^2 S_1 \end{pmatrix}\right), \quad (5)$$

and similarly for other $\gamma = j$ (recall $\mu_1 = \bar{R}_{-1,1}$). The PIP in this model is:

$$p(\gamma = j \mid x) \propto \exp\left(\left(\frac{N\sigma_0^2}{1 + N\sigma_0^2} - \frac{N\sigma_0^2\tau^2}{1 + N\sigma_0^2\tau^2}\right) \frac{Nx_j^2}{2}\right). \quad (6)$$

Note that $\tau^2 = 1/(\nu_0 + 3) \leq 1/3 < 1$, so that $N\sigma_0^2\tau^2 < N\sigma_0^2$, implying

$$\frac{N\sigma_0^2}{1 + N\sigma_0^2} - \frac{N\sigma_0^2\tau^2}{1 + N\sigma_0^2\tau^2} > 0.$$

This quantity tends to $\frac{N\sigma_0^2}{1 + N\sigma_0^2}$ as $\tau^2 \rightarrow 0$ (recovering susie-rss). One interpretation of (6) is that PIP might be more a little more diffuse when we are more uncertain about the in-sample R (which can be a bit funny when we think about it, because if $L = 1$ the LD does not matter and the PIP ideally should be similar to susie-rss. But I guess nothing is perfect, and this really shows how Susie-relax does “relax” the signal.)

However, one can show that this PIP retains some important properties:

1. PIP of SNP j is exponentially proportional to the Z-score square
2. if two SNPs have similar Z-score, they will have the same PIP
3. To compare this PIP with susie-rss, write

$$t = N\sigma_0^2, \quad s = \tau^2.$$

The susie-rss coefficient in the exponent is

$$a_{\text{rss}} = \frac{t}{1 + t},$$

whereas the SER-relax-prime coefficient is

$$a_{\text{prime}} = \frac{t}{1 + t} - \frac{ts}{1 + ts}.$$

Thus

$$0 \leq a_{\text{rss}} - a_{\text{prime}} = \frac{ts}{1 + ts} \leq ts.$$

If $\nu_0 \geq 10$, then $s = \tau^2 = 1/(\nu_0 + 3) \leq 1/13$, so

$$a_{\text{rss}} - a_{\text{prime}} \leq \frac{t}{13 + t} \leq \frac{t}{13}.$$

Equivalently, for any two SNPs j and k ,

$$\left| \log \frac{p_{\text{prime}}(\gamma = j \mid x)}{p_{\text{prime}}(\gamma = k \mid x)} - \log \frac{p_{\text{rss}}(\gamma = j \mid x)}{p_{\text{rss}}(\gamma = k \mid x)} \right| = \frac{N}{2} \frac{ts}{1 + ts} |x_j^2 - x_k^2| \leq \frac{N}{2} \frac{t}{13 + t} |x_j^2 - x_k^2|.$$

Hence, when $t = N\sigma_0^2$ is not too large, $\nu_0 = 10$ already gives only a small flattening of the susie-rss log-odds. The approximation becomes exact as $\nu_0 \rightarrow \infty$.

Posterior distribution of β under SER-relax-prime. Under (5), the covariance no longer depends on β . Therefore, conditional on $\gamma = j$, the likelihood as a function of β is

$$p(x \mid \beta, \gamma = j) \propto \exp\left(N\beta x_j - \frac{N}{2}\beta^2\right),$$

where all remaining terms are constant in β . Combining this likelihood with the prior $\beta \sim \mathcal{N}(0, \sigma_0^2)$ gives

$$p(\beta \mid x, \gamma = j) = \mathcal{N}(m_j, v), \quad (7)$$

$$v = (N + \sigma_0^{-2})^{-1}, \quad (8)$$

$$m_j = vNx_j. \quad (9)$$

Thus the conditional posterior variance is shared across coordinates, and the posterior mean depends on the data only through the active marginal statistic x_j . In particular, $p(\beta \mid x, \gamma = j)$ does not depend on $q = x^T \bar{R}^{-1}x$.

2 susie-relax-prime

How to extend (5) to the sum of a single-effect type of model? Ideally, each of the single-effect models should have the same noise model (like $\mathcal{N}(0, \bar{R}/N)$) to enable IBSS-like inference. I think it is cleanest to introduce a latent variable η_ℓ and an effect contribution vector θ_ℓ for each single-effect component. For $\ell = 1, \dots, L$,

$$\gamma_\ell \sim \text{Unif}\{1, \dots, J\}, \quad \beta_\ell \sim \mathcal{N}(0, \sigma_{0\ell}^2).$$

Conditional on $\gamma_\ell = j$, define

$$\theta_\ell = \beta_\ell \bar{R}_{\cdot j} + \eta_{\ell j},$$

where $\eta_{\ell j, j} = 0$ and

$$\eta_{\ell j, -j} \sim \mathcal{N}(0, \sigma_{0\ell}^2 \tau^2 S_j).$$

The observation model is

$$x \mid \theta_1, \dots, \theta_L \sim \mathcal{N}\left(\sum_{\ell=1}^L \theta_\ell, \frac{\bar{R}}{N}\right).$$

This is a proper generative model. If we integrate out $\eta_{\ell j}$ for a single effect, we recover the SER-relax-prime covariance in (5), with σ_0^2 replaced by $\sigma_{0\ell}^2$.

Variational family and ELBO. Use the mean-field approximation

$$q = \prod_{\ell=1}^L q_\ell(\gamma_\ell, \beta_\ell, \eta_\ell),$$

with

$$q_\ell(\gamma_\ell = j, \beta_\ell, \eta_{\ell j}) = \alpha_{\ell j} q_{\ell j}(\beta_\ell) q_{\ell j}(\eta_{\ell j}).$$

The ELBO is the standard variational objective

$$\begin{aligned}\mathcal{L}(q) = & E_q \log p(x \mid \theta_1, \dots, \theta_L) + \sum_{\ell=1}^L E_q \log p(\gamma_\ell, \beta_\ell, \eta_\ell \mid \gamma_\ell) \\ & - \sum_{\ell=1}^L E_q \log q_\ell(\gamma_\ell, \beta_\ell, \eta_\ell).\end{aligned}\tag{10}$$

Holding all factors except q_ℓ fixed, the optimal coordinate update is

$$\log q_\ell^*(\gamma_\ell, \beta_\ell, \eta_\ell) = E_{-\ell} \log p(x, \{\gamma_k, \beta_k, \eta_k\}_{k=1}^L) + \text{const.}$$

Let

$$\bar{\theta}_k = E_q(\theta_k), \quad r_\ell = x - \sum_{k \neq \ell} \bar{\theta}_k.$$

Then the q_ℓ update is exactly the single-effect SER-relax-prime posterior computed with x replaced by r_ℓ .

Coordinate posterior. For each effect define

$$\kappa_\ell = \frac{N\sigma_{0\ell}^2}{1 + N\sigma_{0\ell}^2}, \quad \rho_\ell = \frac{N\sigma_{0\ell}^2\tau^2}{1 + N\sigma_{0\ell}^2\tau^2}.$$

The coordinate weights are

$$\alpha_{\ell j} = \frac{\exp\left\{\frac{N}{2}(\kappa_\ell - \rho_\ell)r_{\ell j}^2\right\}}{\sum_{k=1}^J \exp\left\{\frac{N}{2}(\kappa_\ell - \rho_\ell)r_{\ell k}^2\right\}}.\tag{11}$$

The common factor involving $q_\ell^{\text{quad}} = r_\ell^T \bar{R}^{-1} r_\ell$ cancels in this normalization. Conditional on $\gamma_\ell = j$,

$$q_{\ell j}(\beta_\ell) = \mathcal{N}(m_{\ell j}, v_\ell),\tag{12}$$

$$v_\ell = (N + \sigma_{0\ell}^{-2})^{-1},\tag{13}$$

$$m_{\ell j} = v_\ell N r_{\ell j} = \kappa_\ell r_{\ell j}.\tag{14}$$

The relaxed part also has a Gaussian posterior. Its nonzero block is

$$q_{\ell j}(\eta_{\ell j, -j}) = \mathcal{N}(\bar{\eta}_{\ell j, -j}, W_{\ell j}),\tag{15}$$

$$\bar{\eta}_{\ell j, -j} = \rho_\ell(r_{\ell, -j} - \mu_j r_{\ell j}),\tag{16}$$

$$W_{\ell j} = \frac{\sigma_{0\ell}^2 \tau^2}{1 + N\sigma_{0\ell}^2 \tau^2} S_j.\tag{17}$$

Equivalently, as a J -vector,

$$\bar{\eta}_{\ell j} = \rho_\ell(r_\ell - r_{\ell j} \bar{R}_{\cdot j}).$$

Posterior mean contribution. The posterior mean contribution of effect ℓ must include both the posterior mean of $\beta_\ell \bar{R}_{\cdot j}$ and the posterior mean of the relaxed component $\eta_{\ell j}$:

$$\begin{aligned}\bar{\theta}_\ell &= \sum_{j=1}^J \alpha_{\ell j} [m_{\ell j} \bar{R}_{\cdot j} + \rho_\ell (r_\ell - r_{\ell j} \bar{R}_{\cdot j})] \\ &= \bar{R}(\alpha_\ell \odot m_\ell) + \rho_\ell r_\ell - \rho_\ell \bar{R}(\alpha_\ell \odot r_\ell).\end{aligned}\tag{18}$$

This correction is what makes the residual update consistent with the variational posterior under the latent contribution model.

ELBO computation. Let $\Omega = \bar{R}^{-1}$ and

$$C_x = -\frac{J}{2} \log(2\pi) - \frac{1}{2} \log \left| \frac{\bar{R}}{N} \right|.$$

For each (ℓ, j) define

$$c_\ell = \frac{\sigma_{0\ell}^2 \tau^2}{1 + N \sigma_{0\ell}^2 \tau^2}, \quad C_{\ell j} = (r_{\ell, -j} - \mu_j r_{\ell j})^T S_j^{-1} (r_{\ell, -j} - \mu_j r_{\ell j}).$$

Using the block inverse identity,

$$C_{\ell j} = r_\ell^T \Omega r_\ell - r_{\ell j}^2,$$

so all $C_{\ell j}$ can be computed from the scalar $r_\ell^T \Omega r_\ell$ and the coordinate-wise residuals. The posterior second moments are

$$B_{\ell j} = E_q(\beta_\ell^2 \mid \gamma_\ell = j) = m_{\ell j}^2 + v_\ell, \tag{19}$$

$$D_{\ell j} = E_q(\eta_{\ell j, -j}^T S_j^{-1} \eta_{\ell j, -j}) = \rho_\ell^2 C_{\ell j} + (J-1)c_\ell. \tag{20}$$

Because $\bar{R}_{\cdot j}^T \Omega = e_j^T$ and $\eta_{\ell j, j} = 0$, the cross term between $\beta_\ell \bar{R}_{\cdot j}$ and $\eta_{\ell j}$ is zero. Hence

$$E_q(\theta_\ell^T \Omega \theta_\ell) = \sum_{j=1}^J \alpha_{\ell j} (B_{\ell j} + D_{\ell j}).$$

Let $\bar{\theta} = \sum_{\ell=1}^L \bar{\theta}_\ell$. A fast way to compute the expected quadratic loss is

$$E_q \left[\left(x - \sum_{\ell=1}^L \theta_\ell \right)^T \Omega \left(x - \sum_{\ell=1}^L \theta_\ell \right) \right] = (x - \bar{\theta})^T \Omega (x - \bar{\theta}) + \sum_{\ell=1}^L V_\ell,$$

where

$$V_\ell = \sum_{j=1}^J \alpha_{\ell j} (B_{\ell j} + D_{\ell j}) - \bar{\theta}_\ell^T \Omega \bar{\theta}_\ell.$$

Thus the expected log-likelihood is

$$E_q \log p(x \mid \theta_1, \dots, \theta_L) = C_x - \frac{N}{2} \left[(x - \bar{\theta})^T \Omega (x - \bar{\theta}) + \sum_{\ell=1}^L V_\ell \right].$$

The remaining terms are most compactly written as KL divergences. Define

$$\text{KL}_{\ell j}^{\beta} = \frac{1}{2} \left[\log \frac{\sigma_{0\ell}^2}{v_{\ell}} + \frac{B_{\ell j}}{\sigma_{0\ell}^2} - 1 \right], \quad (21)$$

and

$$\begin{aligned} \text{KL}_{\ell j}^{\eta} &= \frac{1}{2} \left[(J-1) \log \frac{\sigma_{0\ell}^2 \tau^2}{c_{\ell}} + \frac{D_{\ell j}}{\sigma_{0\ell}^2 \tau^2} - (J-1) \right] \\ &= \frac{1}{2} \left[(J-1) \log(1 + N \sigma_{0\ell}^2 \tau^2) + \frac{D_{\ell j}}{\sigma_{0\ell}^2 \tau^2} - (J-1) \right]. \end{aligned} \quad (22)$$

The determinants $|S_j|$ cancel in $\text{KL}_{\ell j}^{\eta}$ because $W_{\ell j} = c_{\ell} S_j$. The full ELBO is therefore

$$\begin{aligned} \mathcal{L}(q) &= C_x - \frac{N}{2} \left[(x - \bar{\theta})^T \Omega (x - \bar{\theta}) + \sum_{\ell=1}^L V_{\ell} \right] \\ &\quad - \sum_{\ell=1}^L \sum_{j=1}^J \alpha_{\ell j} \left[\log(J \alpha_{\ell j}) + \text{KL}_{\ell j}^{\beta} + \text{KL}_{\ell j}^{\eta} \right]. \end{aligned} \quad (23)$$

This is the expression to use in code. It only requires:

1. the fitted mean $\bar{\theta}$ and one quadratic $(x - \bar{\theta})^T \Omega (x - \bar{\theta})$;
2. for each effect, $\bar{\theta}_{\ell}^T \Omega \bar{\theta}_{\ell}$;
3. for each effect, $r_{\ell}^T \Omega r_{\ell}$, which gives all $C_{\ell j} = r_{\ell}^T \Omega r_{\ell} - r_{\ell j}^2$;
4. coordinate-wise scalar operations for $B_{\ell j}$, $D_{\ell j}$, and the KL terms.

No determinant of S_j is needed, and no dense $W_{\ell j}$ matrix has to be formed.

Variational empirical Bayes updates. Holding the variational distribution fixed, maximize the ELBO over $\sigma_{0\ell}^2$ and τ^2 . Define the effect-level averages

$$B_{\ell} = \sum_{j=1}^J \alpha_{\ell j} B_{\ell j}, \quad D_{\ell} = \sum_{j=1}^J \alpha_{\ell j} D_{\ell j}.$$

The part of the ELBO that depends on $\sigma_{0\ell}^2$ is

$$-\frac{J}{2} \log \sigma_{0\ell}^2 - \frac{1}{2\sigma_{0\ell}^2} \left(B_{\ell} + \frac{D_{\ell}}{\tau^2} \right),$$

so the closed-form VEB update is

$$\hat{\sigma}_{0\ell}^2 = \frac{1}{J} \left(B_{\ell} + \frac{D_{\ell}}{\tau^2} \right).$$

This is the average posterior second moment of the latent effect: one degree of freedom comes from β_ℓ , and $J - 1$ degrees of freedom come from the relaxed LD-column perturbation after rescaling by τ^2 .

Holding $\{\sigma_{0\ell}^2\}_{\ell=1}^L$ fixed, the τ^2 -dependent part of the ELBO is

$$-\frac{L(J-1)}{2} \log \tau^2 - \frac{1}{2\tau^2} \sum_{\ell=1}^L \frac{D_\ell}{\sigma_{0\ell}^2},$$

so

$$\hat{\tau}^2 = \frac{1}{L(J-1)} \sum_{\ell=1}^L \frac{D_\ell}{\sigma_{0\ell}^2}.$$

This update estimates the average standardized size of the relaxed LD-column perturbations. A small value means the posterior does not need much LD relaxation and the method approaches susie-rss; a large value means the data are being explained partly through LD-column uncertainty. In practice these EB updates can be alternated with the variational updates and bounded to avoid degenerate values.

IBSS-like coordinate ascent. The resulting algorithm is coordinate ascent on $\mathcal{L}(q)$:

1. Initialize $\alpha_{\ell j}$, $m_{\ell j}$, v_ℓ , $\bar{\eta}_{\ell j}$, and $\bar{\theta}_\ell$ for $\ell = 1, \dots, L$.

2. For each effect ℓ , form

$$r_\ell = x - \sum_{k \neq \ell} \bar{\theta}_k.$$

3. Compute $v_\ell = (N + \sigma_{0\ell}^{-2})^{-1}$, $m_{\ell j} = v_\ell N r_{\ell j}$, and $\bar{\eta}_{\ell j} = \rho_\ell (r_\ell - r_{\ell j} \bar{R}_{\cdot j})$ for all j .

4. Update the coordinate weights by

$$\alpha_{\ell j} \propto \exp \left\{ \frac{N}{2} (\kappa_\ell - \rho_\ell) r_{\ell j}^2 \right\}.$$

5. Recompute

$$\bar{\theta}_\ell = \bar{R}(\alpha_\ell \odot m_\ell) + \rho_\ell r_\ell - \rho_\ell \bar{R}(\alpha_\ell \odot r_\ell).$$

6. Optionally update $\sigma_{0\ell}^2$ and τ^2 using the VEB formulas, then refresh κ_ℓ , ρ_ℓ , v_ℓ , and $W_{\ell j}$.

7. Sweep over $\ell = 1, \dots, L$ until the ELBO or fitted mean $\sum_{\ell=1}^L \bar{\theta}_\ell$ stabilizes.

Comment on the robustness. We knew from the experiments that under LD mismatch, Susie-RSS's signal $R\bar{b}$ can't fit exactly to the data x (previously denoted by v) and therefore needs more L . This model allow for some deviation from $R\bar{b}$ by additionally using $(\eta_\ell)_{\ell=1}^L$. Looking at the final ELBO (23), the choice of L now is a battle between τ^2 and $\sigma_{0\ell}^2$ in equations (21) and (22): If introduce a little more τ^2 can save us the cost of one more $\sigma_{0\ell}^2$, the model will rather set τ^2 to slightly larger (allowing more deviation) and set that extra $\sigma_{0\ell}^2$ to 0, leading to more conservative behavior.